

# Bubble Detection

Filip Peterek

VŠB-TUO

September 2024

# Dataset

- ▶ Dataset from CAS
- ▶ Manual annotation was necessary
- ▶ Original dataset contains noise, blurry bubbles
- ▶ Only bubbles recognizable by human eye are considered
  - ▶ More than a couple pixels in diameter
  - ▶ Not too blurry

# First Training Set

- ▶ Bounding boxes around bubbles
- ▶ Training set consists of 16 images containing 2766 bubbles
- ▶ Testing set consists of 2 images containing 509 bubbles
- ▶ Used to train RCNNs

## Second Training Set

- ▶ Singular bubbles were extracted
- ▶ Roughly 2500 images of singular bubbles
- ▶ Roughly 2000 images of smidges, noise or clusters of bubbles
- ▶ Used to train Tsetlin Machines

# RCNN Detection

- ▶ Faster RCNNs were used for bubble detection
- ▶ Resnet-50-FPN
- ▶ Transfer learning was applied
- ▶ Has trouble detecting very small bubbles
- ▶ Sometimes misses overlapping bubbles
  - ▶ When both bubbles are very dark
  - ▶ Hardly recognizable even by human eye
- ▶ Couldn't suffice by itself
- ▶ Adjusting contrast improved detection


# Segmentation and Tsetlin Machines

- ▶ To catch what RCNN missed, segmentation is used
- ▶ First, contrast is adjusted
- ▶ Then, thresholding is applied
- ▶ Simple flood fill algorithm selects objects in image
- ▶ Histogram of oriented gradients is computed for each object
- ▶ Tsetlin Machine then determines whether object is a singular bubble

# Segmentation and Tsetlin Machines




# Segmentation and Tsetlin Machines





# Segmentation and Tsetlin Machines




# Combining Results

- ▶ RCNN and Segmentation results are merged together
- ▶ Intersection over union is computed
- ▶ If IOU for two objects passes a certain threshold, we consider them the same object
- ▶ Thus, duplicates are filtered
- ▶ Segmentation results are preferred over RCNN results
  - ▶ Prevents margins around bubbles


# Results

- ▶ Around 76% accuracy
- ▶ Six false positives
- ▶ Detection using Resnet takes 236 ms
- ▶ Recognition using Tsetlin Machines takes 79 ms
  - ▶ Includes HOG computation
- ▶ In total, processing one image takes 1506 ms
- ▶ Mainly the result of an inefficient flood fill algorithm
  - ▶ 990 ms
- ▶ Written in Python, can be optimized

# Results




# Results



# Results



# Results




# Challenges

- ▶ Very dark transitions




# Challenges



# Challenges

- ▶ Very dark transitions
- ▶ Blurry bubbles


# Challenges



# Challenges

- ▶ Very dark transitions
- ▶ Blurry bubbles
- ▶ Bright spots where light reflects
  - ▶ Breaks thresholding

# Challenges



# Failed Experiments

- ▶ Local Binary Pattern
- ▶ Background subtraction
- ▶ Convolutional Tsetlin Machine